



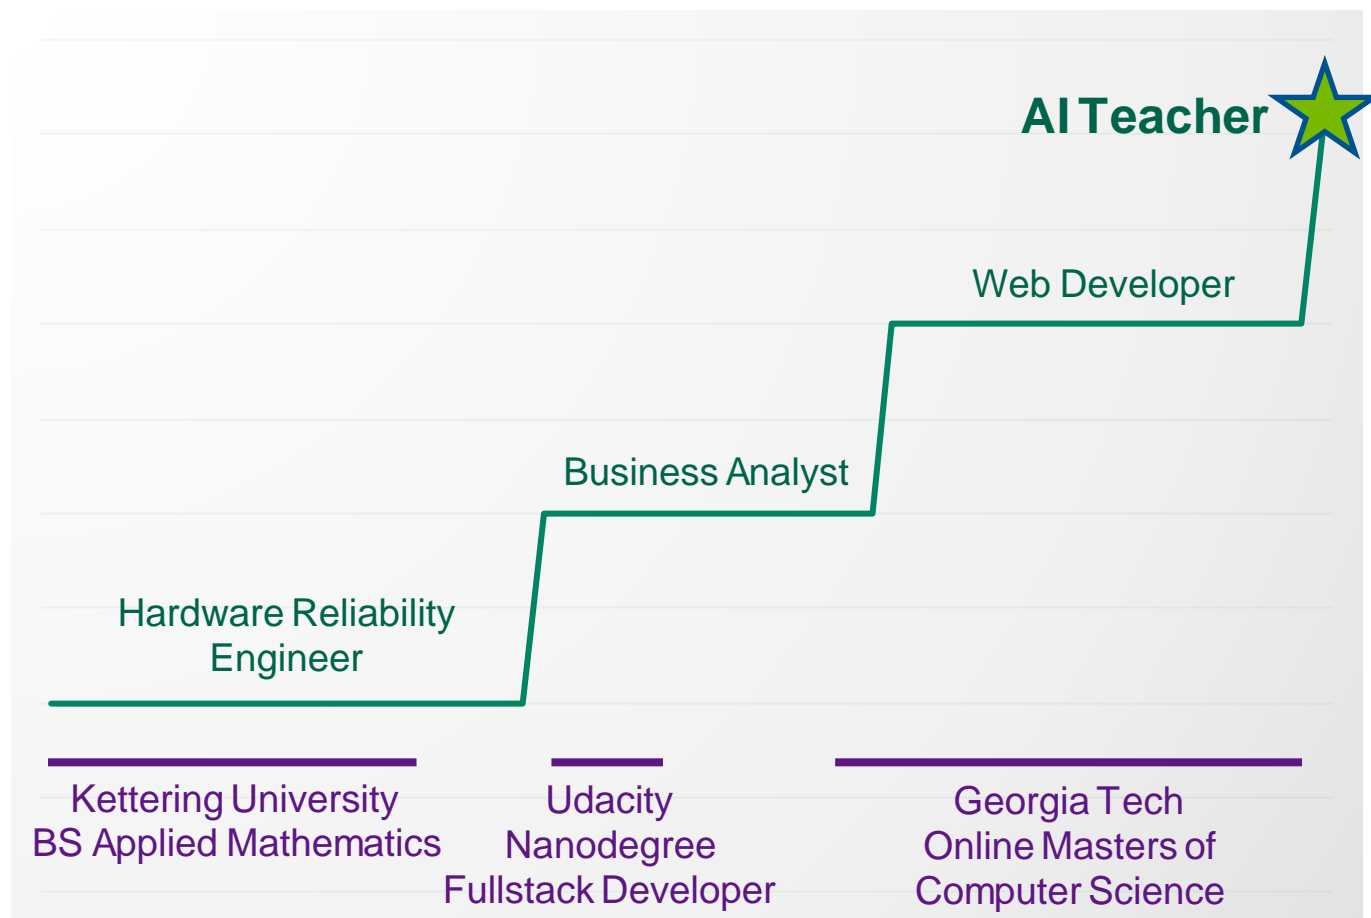
BECOMING A DATA SCIENTIST

Danielle Detering, NVIDIA Deep Learning Institute



Danielle Detering
DLI Content Developer

MY DATA SCIENCE JOURNEY





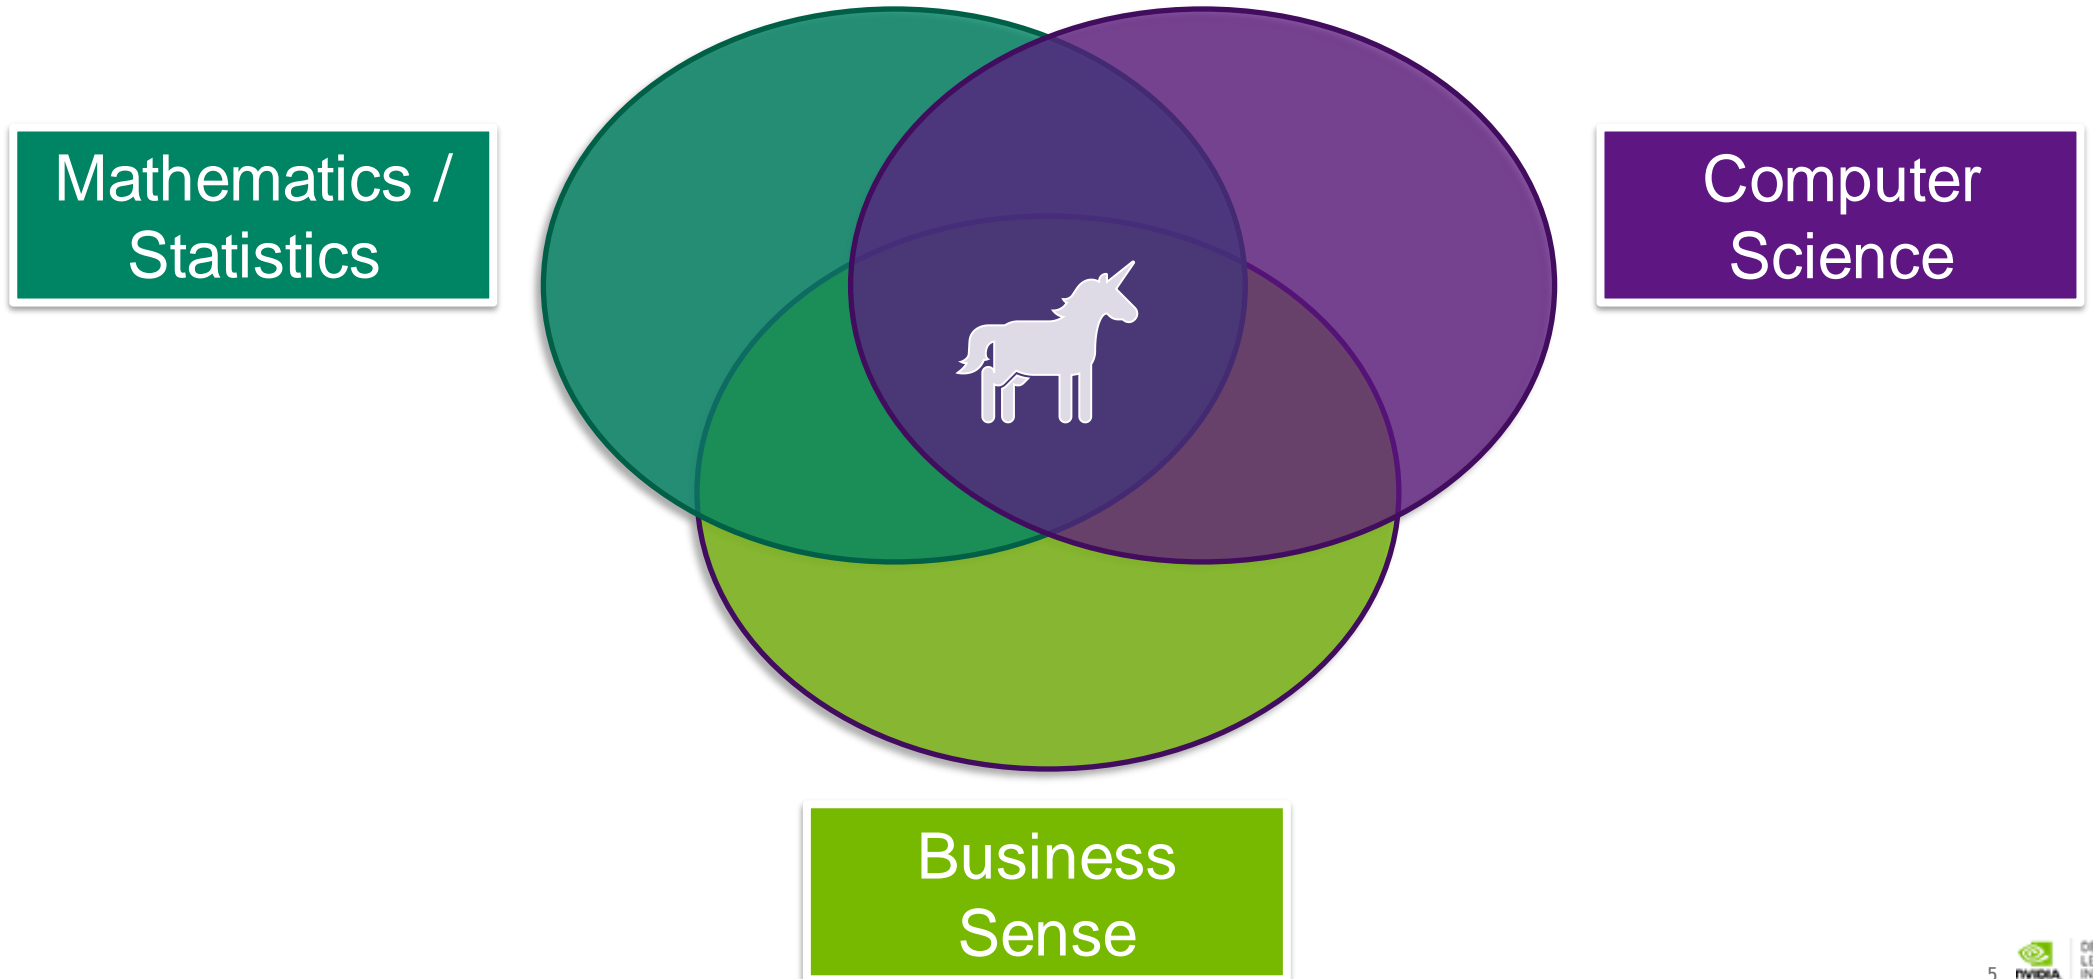
DEEP
LEARNING
INSTITUTE

Training developers, data scientists, researchers and IT professionals how to solve their most challenging problems

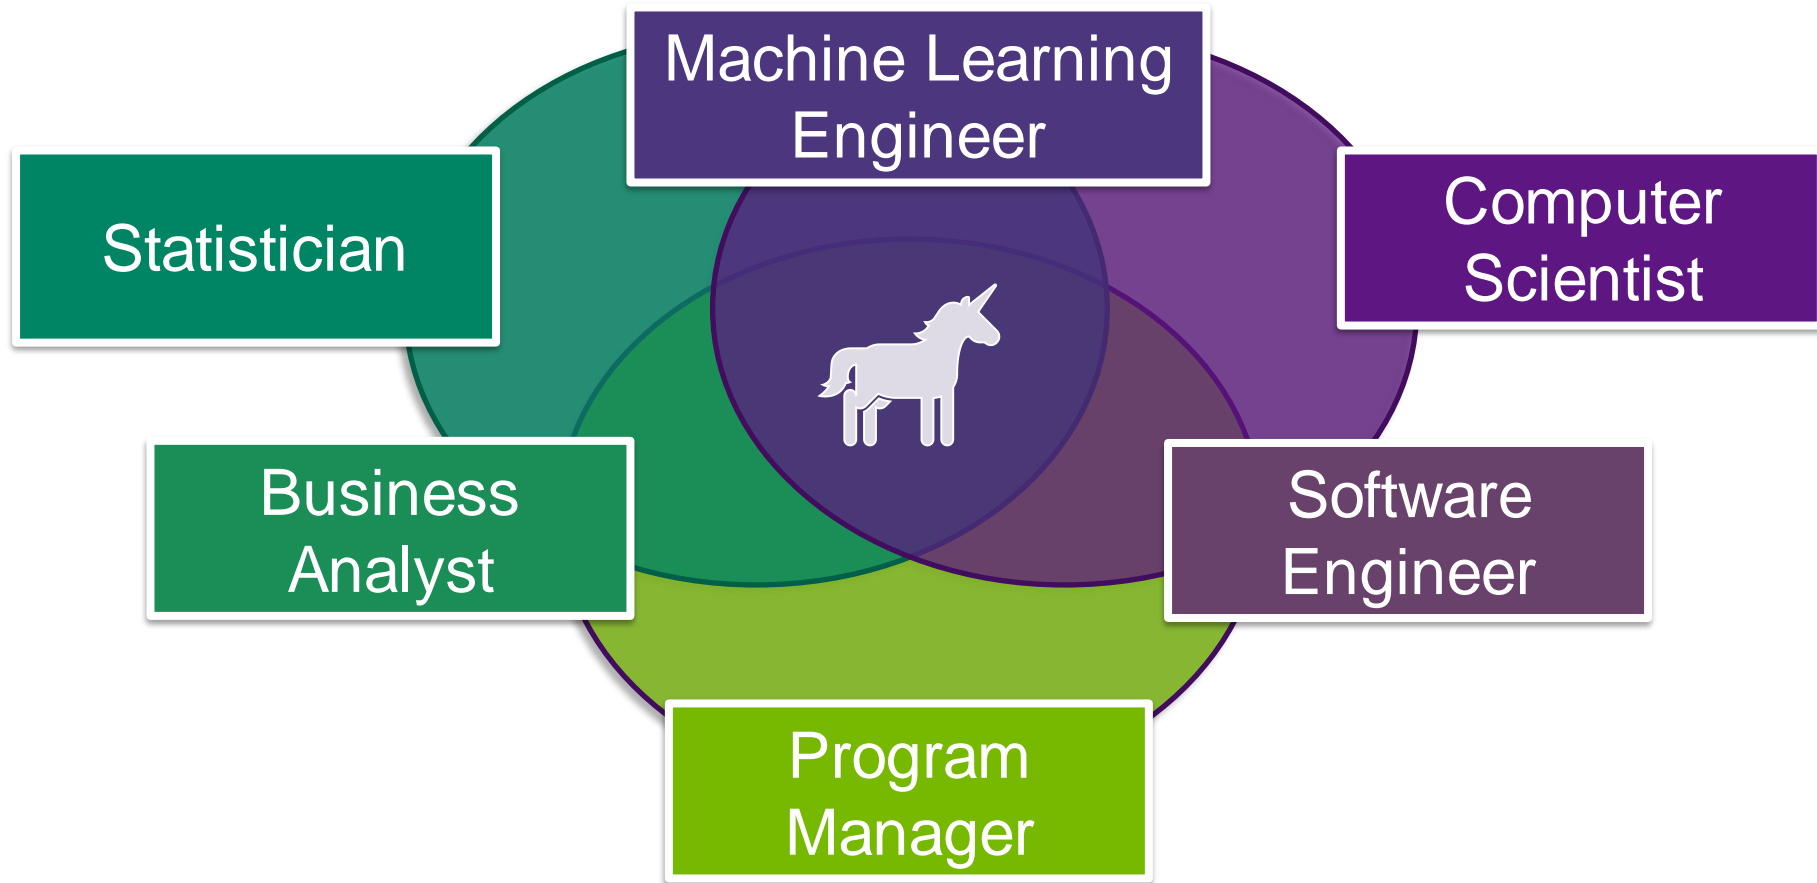
The background is a dark blue gradient with a complex network of thin, light green lines crisscrossing across the frame. At various points where these lines intersect, there are small, bright green circular dots. Some of these dots have a soft, out-of-focus glow around them. The overall effect is one of a digital or neural network, suggesting connectivity and technology.

KEYS TO SUCCESS

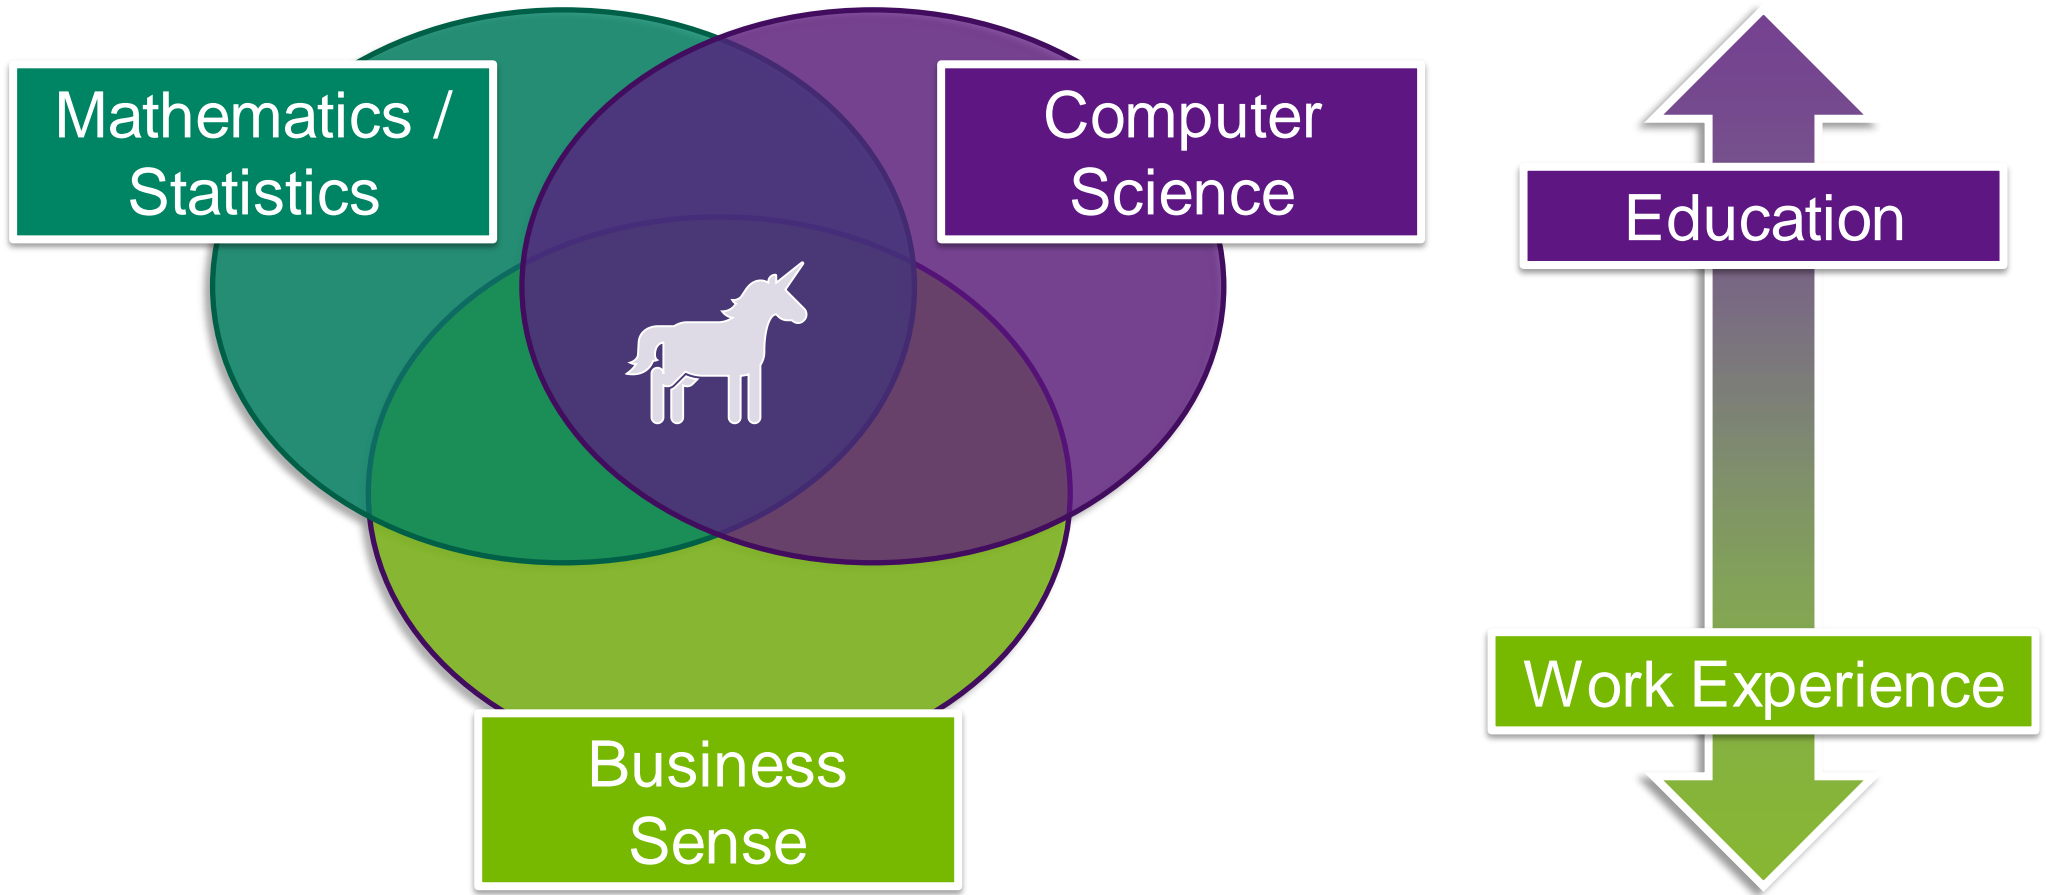
DATA SCIENCE “UNICORNS”



DATA SCIENCE “UNICORNS”



DATA SCIENCE “UNICORNS”



UNIVERSITY VS WORKSHOPS

University “Coarse” Adjustment


- Takes many weeks
- Covers concepts in great depth



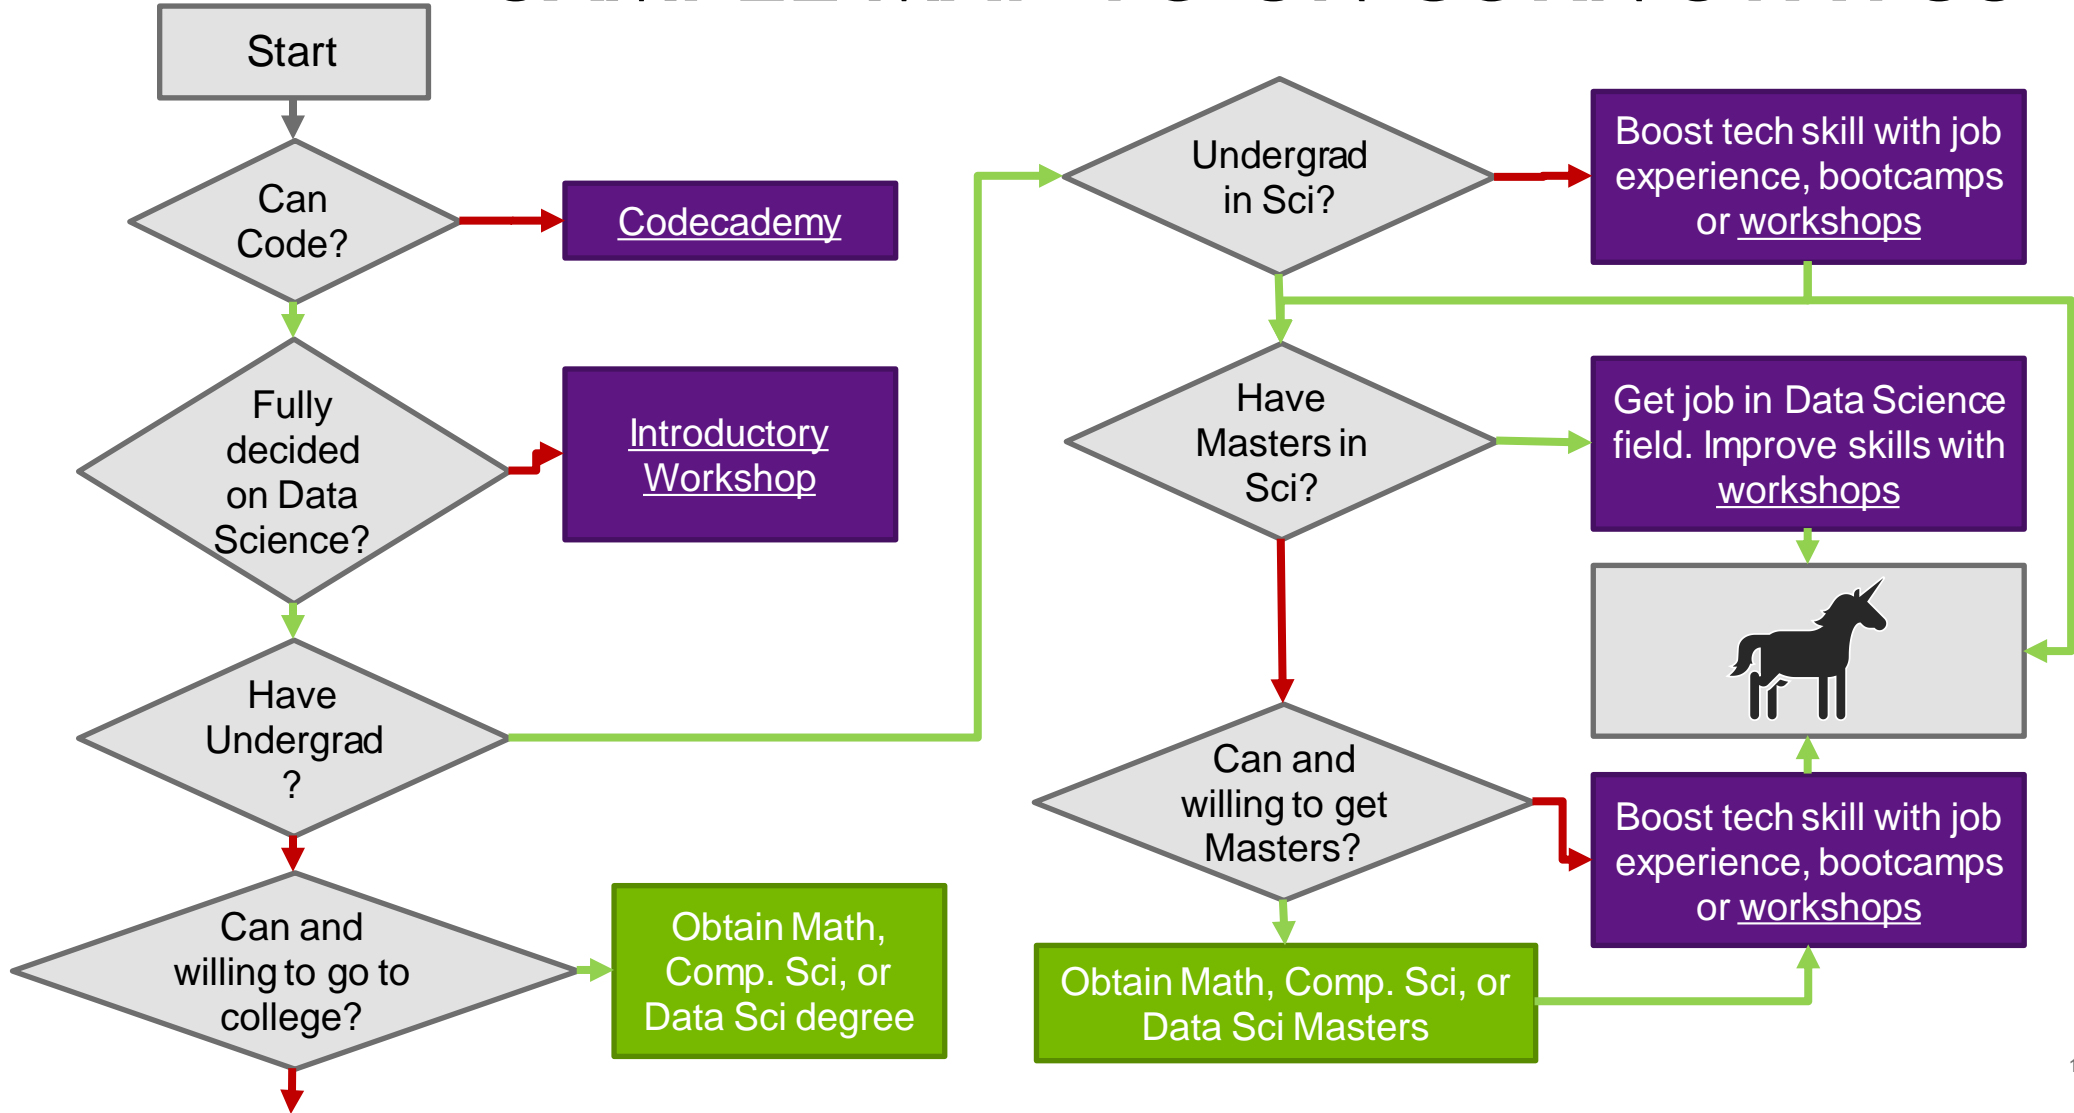
Workshop “Fine” Adjustment

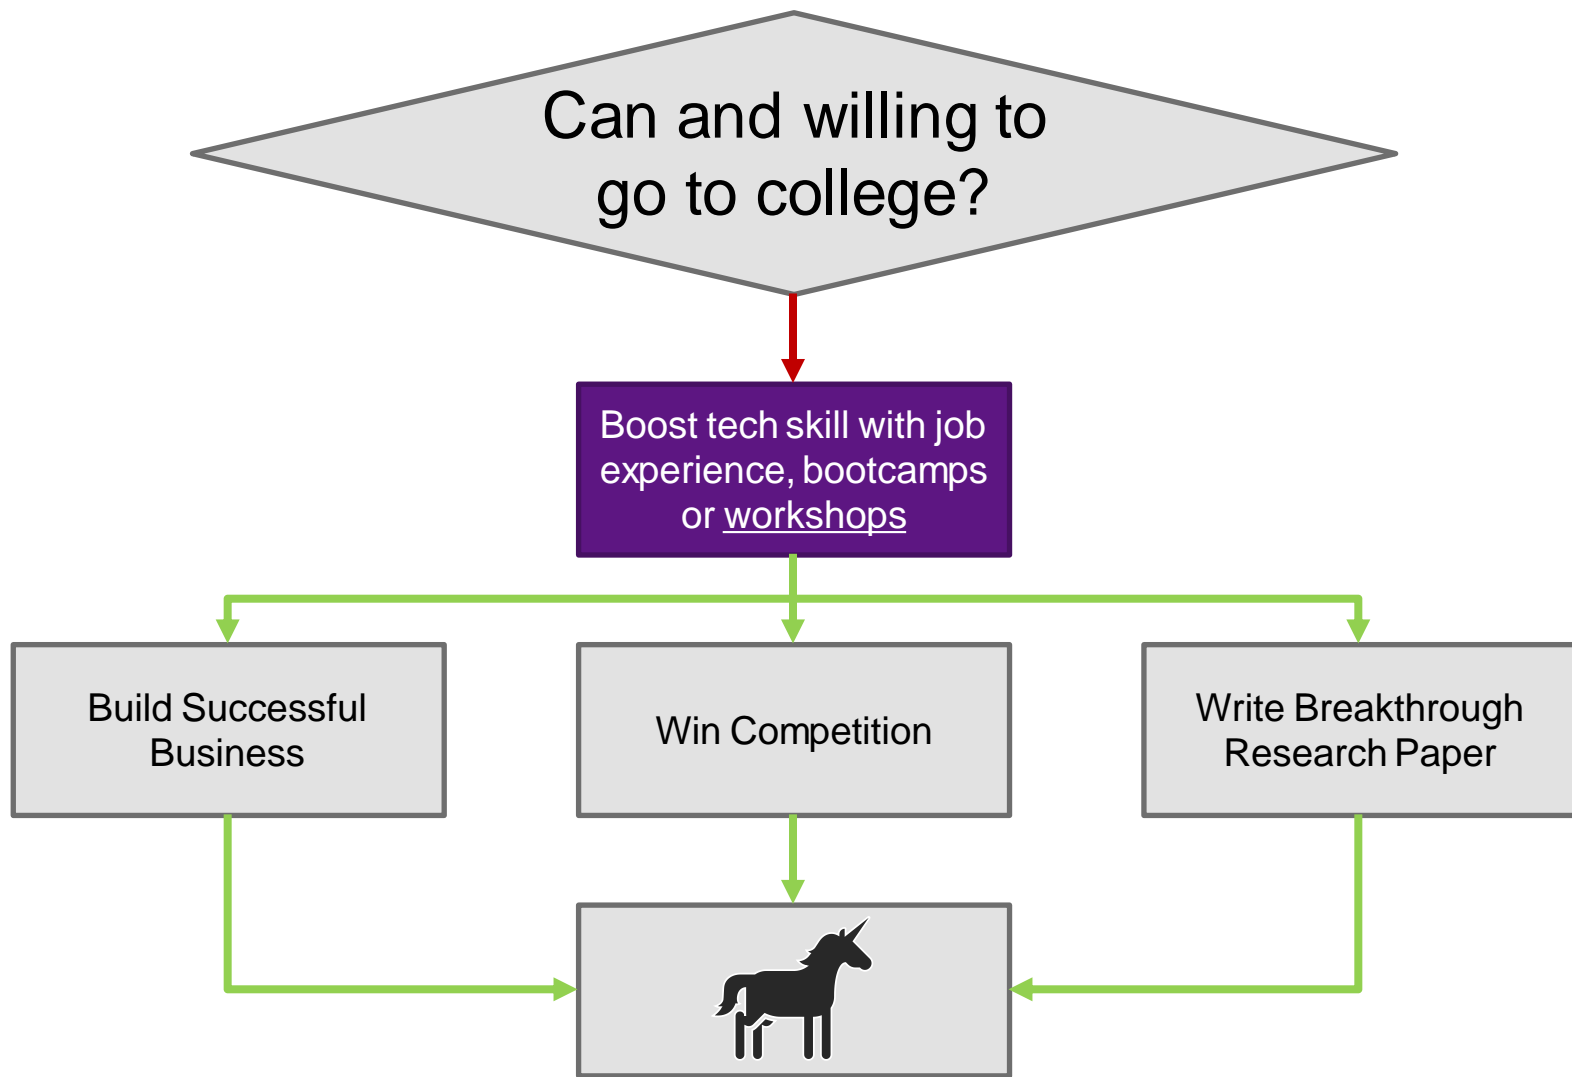
- Takes a few hours or days
- Covers cutting edge topics close to industry

UNIVERSITY VS WORKSHOPS



SAMPLE MAP TO UNICORN STATUS





TIPS FOR NAVIGATING



Find an “Accountability Partner”



Find a goal that inspires you


Teach AI to play games?
Teach an AI to sing?
Teach a robot dog to fetch the paper?

The background is a dark blue gradient with a complex network of thin, glowing green lines. These lines connect various points, some of which are larger, bright green circular nodes. The overall effect is reminiscent of a neural network or a data visualization of complex connections.


NVIDIA DEEP LEARNING INSTITUTE

RICH CONTENT PORTFOLIO

Fundamentals and advanced hands-on training in key technologies and application domains



Deep Learning
Fundamentals



Accelerated Computing
Fundamentals



Accelerated Data Science
Fundamentals



Intro to AI in the Data
Center




AI for Anomaly Detection



AI for Autonomous Vehicles



AI for
Digital Content Creation



AI for Healthcare



AI for Industrial Inspection



AI for
Intelligent Video Analytics



AI for
Predictive Maintenance



AI for Robotics

WHAT DOES DLI OFFER?

SELF-PACED ONLINE

- Get started anywhere, any time with access to a GPU-accelerated workstation in the cloud
- Full-day courses (8 hrs) are \$90
- 2-4 hour courses are \$30
- Bulk pricing is also available

Get started www.nvidia.com/dli

INSTRUCTOR-LED WORKSHOP

- Full-day workshops onsite at your location or remote, delivered by DLI certified instructors
- MSRP: \$10K/day for up to 20 attendees (EDU pricing available)
- Request through your account manager
- Public workshop schedule [here](#).

ENTERPRISE SOLUTIONS

- End-to-end training solution with executive briefings, enterprise-level reporting, and a mix of onsite and online training
- Pricing varies
- Request through your account manager

INSTRUCTOR- LED WORKSHOPS

In-person or Remote

AGENDA

Introduction (45 mins)

Break (15 mins)

Training task #1 (1:20 mins)

Lunch (60 mins)

Training task #2 (1:20 mins)

Break (15 mins)

Training task #3 (1:20 mins)

Summary/Q&A (15 mins)

BENEFITS OF WORKSHOP

Get guidance from DLI Certified Instructors while working through material

Hear from experts during the introductory lecture

Collaborate with and learn from peers

Access fully-configured, GPU-accelerated workstations in the cloud

Earn a certificate of competency in course subject matter

ONLINE TRAINING

BENEFITS OF SELF-PACED TRAINING

Take training any time, anywhere. All you need is a laptop and Internet connection.

Access fully configured, GPU-accelerated workstations in the cloud for hands-on training.

Learn at your own pace with hands-on exercises and videos.

Earn a certificate of competency in 8-hr course subject matter (digital badge of completion for IT training.)

Customers can purchase bulk codes for online courses

DLI UNIVERSITY TRAINING

Learn more at
www.nvidia.com/dli

UNIVERSITY AMBASSADOR PROGRAM

- Qualified faculty and researchers can get certified to teach DLI workshops to their students at no cost.
- Hundreds of universities certified around the world, including:



TEACHING KITS

- Qualified university educators can download courseware across deep learning, accelerated computing, and robotics.
- Kits include lecture materials, GPU cloud resources, access to self-paced DLI courses, and more.

OFFERINGS OVERVIEW

Traditional College



Nanodegree



- [Udacity Self-Driving Car Nanodegree](#)
- [Udacity Robotics Nanodegree](#)
- [Udacity Computer Vision Nanodegree](#)

Workshops

- [Fundamental of Deep Learning for Computer Vision](#)
- [High Performance Computing with Containers](#)
- [And more!](#)

“Coarse” Adjustment

“Fine” Adjustment



ADVANCE YOUR DEEP LEARNING KNOWLEDGE AT GTC

The world's most important event for GPU developers

WWW.GPUTECHCONF.COM




ACCELERATING AI WORKFLOWS WITH NGC




A TYPICAL AI WORKFLOW

Training and Inference

TRAINING



INFERENCE



DIFFERENT ROLES. SAME GOALS.

Driving Productivity and Faster Time-to-Solutions

Data Engineer



Gather, organize, manage data to build models

Data Scientists and Researchers



Model convergence & accuracy

Developers



Integrate models in application software

DevOps



Deploy and maintain ML systems in production reliably and efficiently

Sys Admins



Support internal and external customers

NVIDIA NGC | ACCELERATING TIME TO SOLUTION


SIMPLIFY DEVELOPMENT

NGC Catalog

Containers



Models



Helm Charts



Build | Customize | Integrate

MANAGE & SECURE APPLICATIONS

NGC Private Registry

Access
Control



Scanning, Signing,
Encryption



Lifecycle
Management



Manage | Secure | Share

DEPLOY ANYWHERE

NGC-Ready

On Premises



Cloud



Edge



ENABLING PORTABILITY WITH NGC CONTAINERS

Extensive

- Diverse range of workloads and industry specific use cases

Optimized

- DL containers updated monthly
- Packed with latest features and superior performance

Secure & Reliable

- Scanned for vulnerabilities and crypto
- Tested on workstations, servers, & cloud instances

Scalable

- Supports multi-GPU & multi-node systems

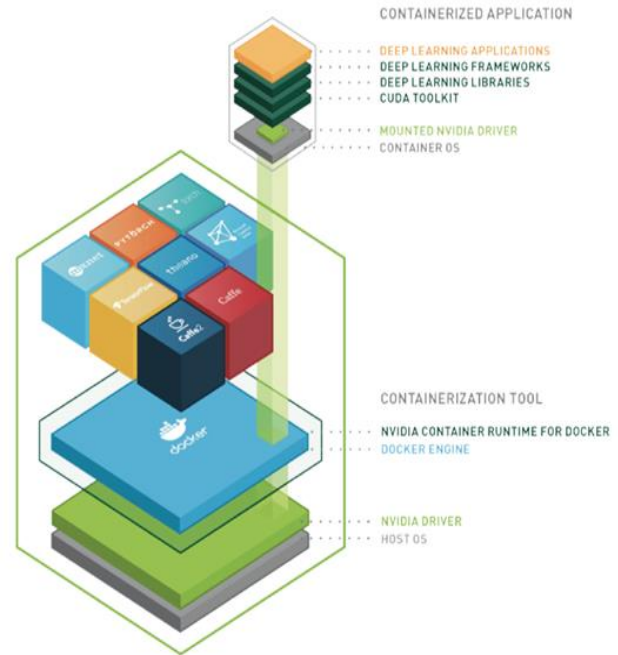
Designed for Enterprise & HPC

- Supports Docker, Singularity & other runtimes

Run Anywhere

- Bare metal, VMs, Kubernetes
- x86, ARM, POWER
- Multi-cloud, on-prem, hybrid, edge

NGC Deep Learning Containers



CONVERSATIONAL AI



JARVIS

HEALTHCARE



CLARA

SMART CITIES



DEEPSTREAM &
SMART PARKING

TELECOM



AERIAL

AUTONOMOUS DRIVING



DRIVE

ROBOTICS



ISAAC

HPC



HPC SDK

[Learn more about NGC Containers](#)

CONTINUOUS PERFORMANCE IMPROVEMENT

Developers' Software Optimizations Deliver Better Performance on the Same Hardware

Monthly DL Framework Updates & Stack Optimizations Drive Performance

cuDNN - Highly tuned standard training routines

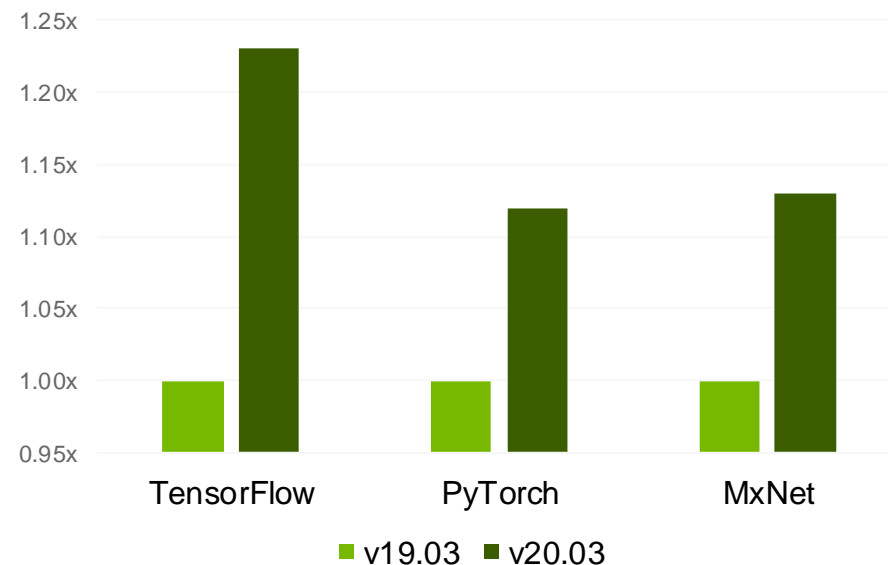
cuBLAS - Highly tuned matrix multiplication

DALI - Moves compute intensive pre-processing to GPUs

NCCL - Faster training across multi-GPU architecture

Framework - Latest versions w/ newest features and superior perf

PERFORMANCE GAINS ACROSS FRAMEWORKS



512 Batch Size for TF & PyT, 256 Batch size for MxNet | ResNet-50 Training v1.5 | 16x V100 | DGX-2

PRE-TRAINED MODELS AND MODEL SCRIPTS

Build AI Faster


PRE-TRAINED MODELS

- Popular AI tasks - ASR, NLU, TTS, RecSys, CV, etc
- Industry specific models - Medical imaging, public safety
- Customize with your data and transfer learning
- Integrate into existing workflows with SDKs

MODEL SCRIPTS

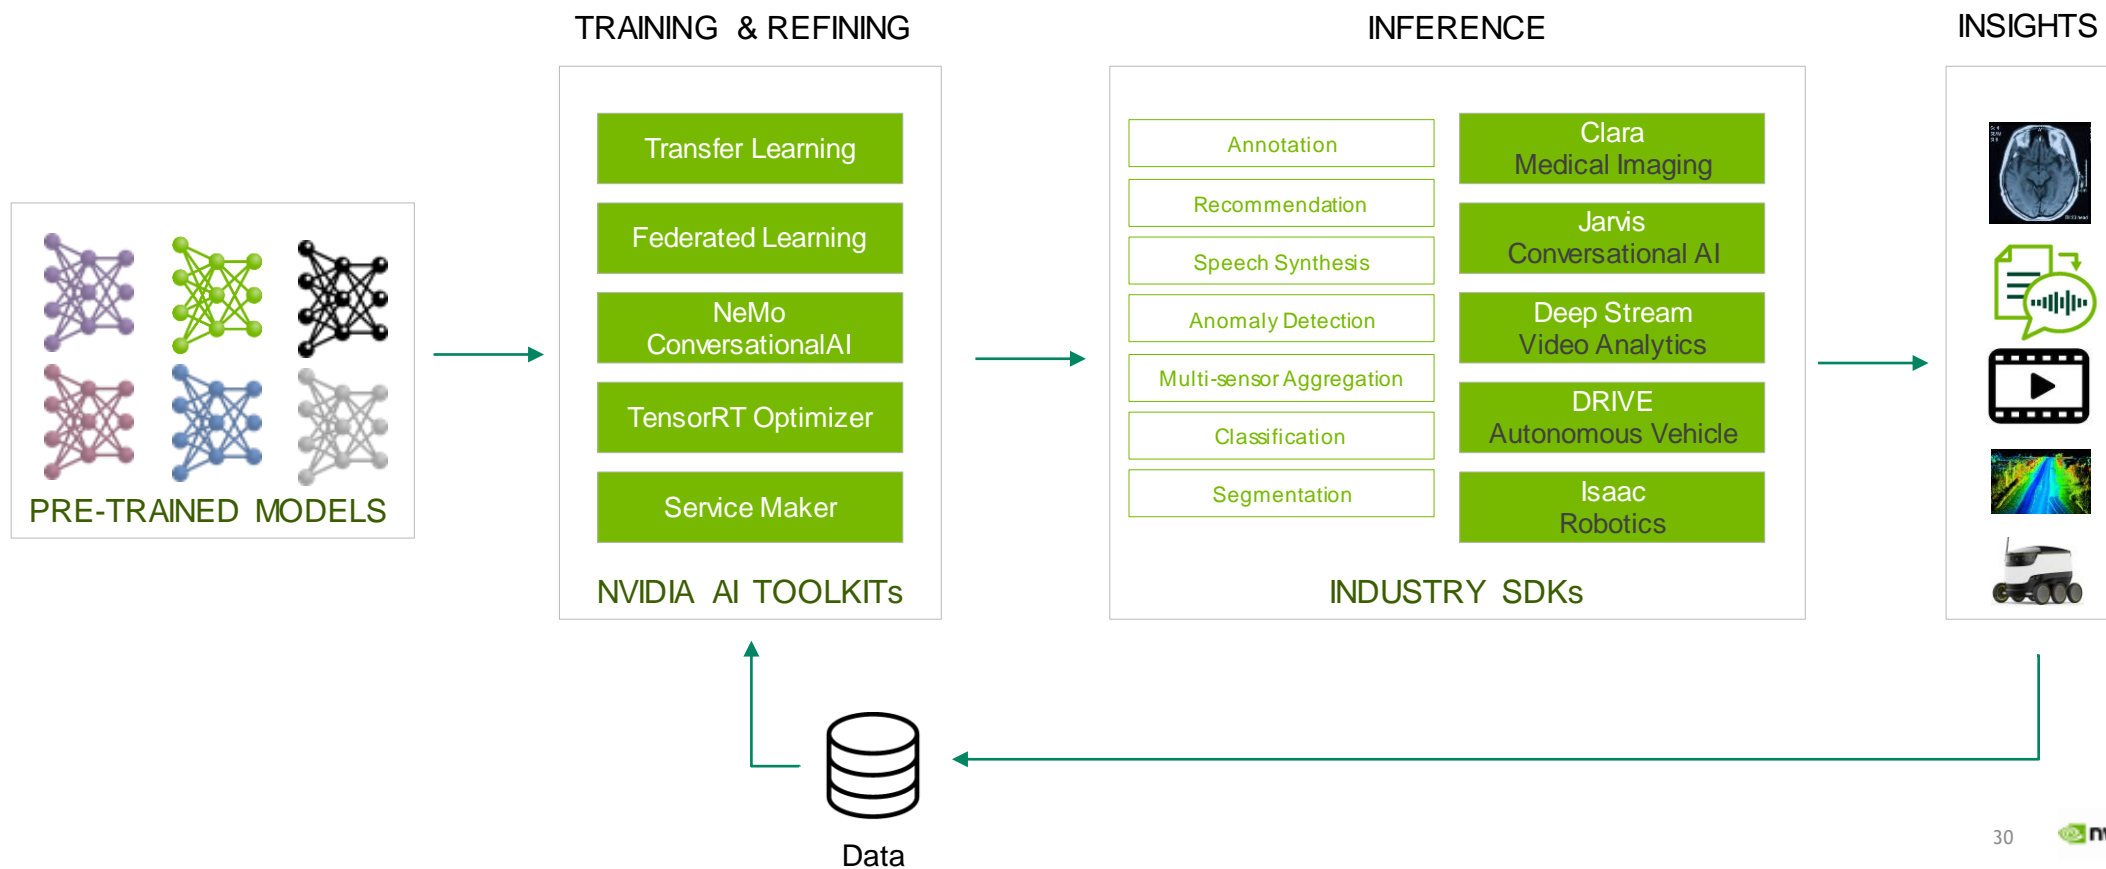
- Reference neural network architectures across all domains and popular frameworks with latest SOTA
- Code samples show you how to deploy or build your models

FASTEST INFERENCE WITH GPUS



AI TOOLKITs AND SDKs

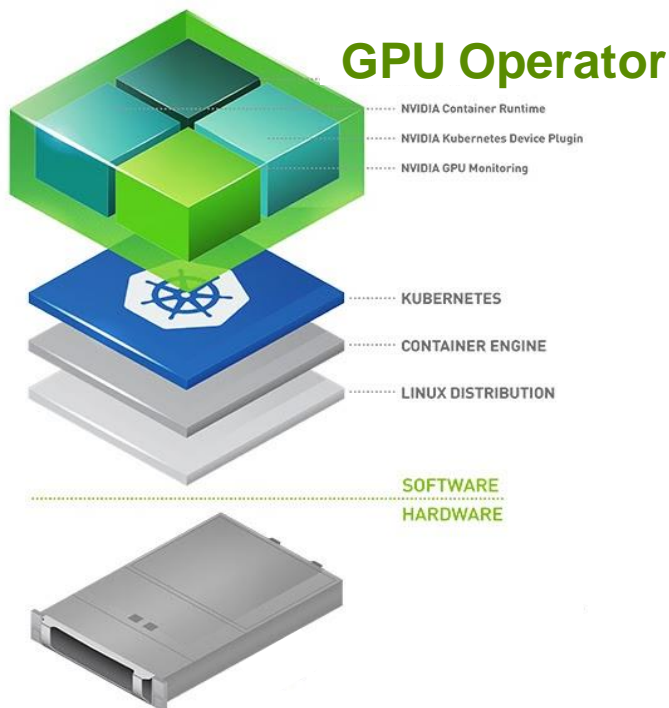
Build AI Even Faster



CLOUD NATIVE DEPLOYMENT APPROACH

Enabling Consistent Deployment Across Platforms with Kubernetes

NVIDIA EGX Stack



kubernetes

Container Orchestration
Automated container deployment including
self-healing



Helm is NVIDIA's recommended package manager for Kubernetes that allows you to more easily configure, deploy and update applications on Kubernetes

NGC PRIVATE REGISTRY

Secure, Cloud-hosted Platform to Collaborate and Build Faster

Access Control

- User roles allow you to discover, share and control your assets


Secure

- Encryption at rest and in transit
- Automated Security Scanning to protect your container images against a comprehensive list of CVEs, crypto keys
- Container signing to ensure the software comes from a trusted source and hasn't been tampered with

Lifecycle Management

- Build and share your artifacts with cross-functional teams
- Version control system designed for models
- Use NGC's flexible API, CLI and Web UI to build workflows

Private Registry Documentation



WHY NGC

Optimized for Enterprise Needs

CURATED SOFTWARE

FASTER TIME TO SOLUTION

- Built & maintained by experts
- Covers popular applications & use cases
- Supercharged with latest features

SUPERIOR PERFORMANCE

RUN LARGER MODELS/SIMULATION

- AI s/w constantly optimized
- Instantly access latest features and highest performance
- Winner of MLPerf competition

TESTED ACROSS PLATFORMS

RELIABLE SOFTWARE

- Supports multi-GPU & multi-node systems
- Passes stringent security scans
- Cloud | Edge | On-prem | Workstations

ENTERPRISE-GRADE SUPPORT

DEPLOY WITH CONFIDENCE

- Access to NVIDIA AI experts
- Faster time-to-solution
- Minimizes system downtimes

BUILD AI FASTER. DEPLOY ANYWHERE WITH NGC

ngc.nvidia.com

The screenshot displays the NVIDIA NGC Accelerated Software interface. The top navigation bar includes the NVIDIA NGC logo, the text "ACCELERATED SOFTWARE", and a "SIGN IN" link. Below the navigation bar, there are tabs for "ALL CONTENT TYPES", "CONTAINERS", "MODELS", "MODEL SCRIPTS", and "HELM CHARTS". The "ALL CONTENT TYPES" tab is selected, and a search bar is visible. The main content area shows a grid of software packages, each with a logo, title, description, version, and build date. The packages are organized into three rows and eight columns.

Package Name	Version	Build Date
ParaView IndexX	5.7.0-egl-pvw	02/21/20
Isaac Sim	2020.1 preview	02/21/20
Conundrum Aircraft-En...	latest	02/20/20
Unsupervised Latent La...	1	02/20/20
V-Net Medical for Tenso...		02/20/20
BERT for PyTorch		02/20/20
BERT for TensorFlow		02/20/20
Transformer-XL for PyTo...		02/20/20
NCF for PyTorch		02/20/20
OpenACC Training Mater...	20.1.1	02/20/20
Clara-Train-SDK	v2.0	02/18/20
Transformer-BIG-en-de...	1	02/13/20
JasperNetDr 10x5 for N...	1	02/13/20
BertLargeUncasedForN...	4	02/11/20
BertBaseCasedForNeMo	3	02/11/20
BertBaseUncasedForNe...	2	02/11/20
RAPIDS	cuda10.1-runtime-ubuntu18.04	02/06/20
NVCaffe	20.01-py3	02/04/20
DIGITS	20.01-tensorflow-py3	01/27/20
MXNet	20.01-py3	01/27/20
Kaldi	20.01-py3	01/27/20
TensorRT Inference Serv...	20.01-py3	01/27/20
PyTorch	20.01-tf1-py2	01/27/20
TensorFlow	20.01-tf1-py2	01/27/20

WHAT'S NEXT

GET THE BASICS

Watch [“Deep Learning and Beyond”](#)

Listen to the [NVIDIA AI Podcast](#)

Review [examples of AI in action](#)

LEARN WITH DLI

Take a self-paced online training at
www.nvidia.com/dli

Request an onsite or remote workshop through
your account manager

JOIN OUR COMMUNITY

Sign up for the NVIDIA Developer Program at
<https://developer.nvidia.com/join>

NGC Catalog

Check out the available containers and resources
in the [NGC Catalog](#)

See the [docs](#) to get started



nvidia®

www.nvidia.com/dli
www.nvidia.com/gpu-cloud/